



Exercise 5

Image Classification

▶ Cifar 10 dataset

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



<https://www.cs.toronto.edu/~kriz/cifar.html>

Image Classification


▶ Cifar 10 dataset



airplane
automobile
bird
cat
deer
dog
frog
horse
ship
truck

Coding 전에 생각해 볼 것

- ▶ 입력 데이터 : `image(C , H , W)`
- ▶ 출력 데이터 : `class number`
- ▶ Optimizer: ?
- ▶ Loss function: `cross entropy`



준비 단계 1

▶ 입, 출력 데이터

```
transform_CIFAR10 = transforms.Compose([
    transforms.ToTensor()
])
train_loader = torch.utils.data.DataLoader(
    datasets.CIFAR10(
        root      = './data_CIFAR10',
        train      = True,
        download   = True,
        transform  = transform_CIFAR10),
    batch_size=BATCH_SIZE,
    shuffle=True
)
test_loader = torch.utils.data.DataLoader(
    datasets.CIFAR10(
        root      = './data_CIFAR10',
        train      = False,
        download   = True,
        transform  = transform_CIFAR10),
    batch_size=BATCH_SIZE,
    shuffle=True
)
```

▶ Model

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 30, 30]	224
Conv2d-2	[-1, 16, 28, 28]	1,168
Conv2d-3	[-1, 24, 26, 26]	3,480
Dropout2d-4	[-1, 24, 26, 26]	0
Linear-5	[-1, 128]	519,296
Linear-6	[-1, 10]	1,290

Total params: 525,458

Trainable params: 525,458

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 0.40

Params size (MB): 2.00

Estimated Total Size (MB): 2.42

training and validation

```
: def train(model, train_loader, optimizer, epoch):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(DEVICE), target.to(DEVICE)
        optimizer.zero_grad()
        output = model(data)
        loss = F.cross_entropy(output, target)
        loss.backward()
        optimizer.step()

        if batch_idx % 200 == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
```

```
: def evaluate(model, test_loader):
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(DEVICE), target.to(DEVICE)
            output = model(data)
            test_loss += F.cross_entropy(output, target,
                                         reduction='sum').item()
            pred = output.max(1, keepdim=True)[1]
            correct += pred.eq(target.view_as(pred)).sum().item()

    test_loss /= len(test_loader.dataset)
    test_accuracy = 100. * correct / len(test_loader.dataset)
    return test_loss, test_accuracy
```

```
for epoch in range(1, EPOCHS + 1):
    train(model, train_loader, optimizer, epoch)
    test_loss, test_accuracy = evaluate(model, test_loader)

    print('[{}] Test Loss: {:.4f}, Accuracy: {:.2f}%'.format(
        epoch, test_loss, test_accuracy))
```

Question and Answer